

# Gradient Descent

---

Hoàng Nam Dũng

Khoa Toán - Cơ - Tin học, Đại học Khoa học Tự nhiên, Đại học Quốc gia Hà Nội

## Gradient descent

Consider unconstrained, smooth convex optimization

$$\min_x f(x)$$

with convex and differentiable function  $f : \mathbb{R}^n \rightarrow \mathbb{R}$ . Denote the optimal value by  $f^* = \min_x f(x)$  and a solution by  $x^*$ .

## Gradient descent

Consider unconstrained, smooth convex optimization


$$\min_x f(x)$$


with convex and differentiable function  $f : \mathbb{R}^n \rightarrow \mathbb{R}$ . Denote the optimal value by  $f^* = \min_x f(x)$  and a solution by  $x^*$ .

**Gradient descent:** choose initial point  $x^{(0)} \in \mathbb{R}^n$ , repeat:

$$x^{(k)} = x^{(k-1)} - t_k \cdot \nabla f(x^{(k-1)}), \quad k = 1, 2, 3, \dots$$

Stop at some point.





## Gradient descent interpretation

At each iteration, consider the expansion

$$f(y) \approx f(x) + \nabla f(x)^T (y - x) + \frac{1}{2t} \|y - x\|_2^2.$$


**Quadratic approximation**, replacing usual Hessian  $\nabla^2 f(x)$  by  $\frac{1}{t} I$ .

$$\begin{array}{ll} f(x) + \nabla f(x)^T (y - x) & \text{linear approximation to } f \\ \frac{1}{2t} \|y - x\|_2^2 & \text{proximity term to } x, \text{ with weight } 1/2t \end{array}$$

Choose next point  $y = x^+$  to minimize quadratic approximation

$$x^+ = x - t \nabla f(x).$$

## Gradient descent interpretation



Blue point is  $x$ , red point is

$$x^* = \operatorname{argmin}_y f(x) + \nabla f(x)^T(y - x) + \frac{1}{2t} \|y - x\|_2^2$$


# Outline

- ▶ How to choose step sizes
- ▶ Convergence analysis
- ▶ Nonconvex functions
- ▶ Gradient boosting

## Fixed step size


Simply take  $t_k = t$  for all  $k = 1, 2, 3, \dots$ , can **diverge** if  $t$  is too big.

Consider  $f(x) = (10x_1^2 + x_2^2)/2$ , gradient descent after 8 steps:




## Fixed step size

Can be **slow** if  $t$  is too small. Same example, gradient descent after 100 steps:



## Fixed step size

Converges nicely when  $t$  is “just right”. Same example, 40 steps:



Convergence analysis later will give us a precise idea of “just right”.

## Backtracking line search

One way to adaptively choose the step size is to use backtracking line search:

- ▶ First fix parameters  $0 < \beta < 1$  and  $0 < \alpha \leq 1/2$ .
- ▶ At each iteration, start with  $t = t_{\text{init}}$ , and while


$$f(x - t \nabla f(x)) > f(x) - \alpha t \|\nabla f(x)\|_2^2$$

shrink  $t = \beta t$ . Else perform gradient descent update

$$x^+ = x - t \nabla f(x).$$

Simple and tends to work well in practice (further simplification: just take  $\alpha = 1/2$ ).


## Backtracking interpretation



For us  $\Delta x = -\nabla f(x)$

## Backtracking line search

Setting  $\alpha = \beta = 0.5$ , backtracking picks up roughly the **right step size** (12 outer steps, 40 steps total).



## Exact line search

We could also choose step to do the best we can along direction of negative gradient, called **exact line search**:

$$t = \operatorname{argmin}_{s \geq 0} f(x - s \nabla f(x)).$$

Usually not possible to do this minimization exactly.

Approximations to exact line search are typically not as efficient as backtracking and it's typically not worth it.

## Convergence analysis

Assume that  $f : \mathbb{R}^n \rightarrow \mathbb{R}$  convex and differentiable and additionally

$$\|\nabla f(x) - \nabla f(y)\|_2 \leq L \|x - y\|_2 \quad \text{for any } x, y$$

i.e.,  $\nabla f$  is Lipschitz continuous with constant  $L > 0$ .

### Theorem

*Gradient descent with fixed step size  $t \leq 1/L$  satisfies*

$$f(x^{(k)}) - f^* \leq \frac{1}{2tk} \|x^{(0)} - x^*\|_2^2$$

*and same result holds for backtracking with  $t$  replaced by  $\beta/L$ .*

## Convergence analysis

Assume that  $f : \mathbb{R}^n \rightarrow \mathbb{R}$  convex and differentiable and additionally

$$\|\nabla f(x) - \nabla f(y)\|_2 \leq L \|x - y\|_2 \quad \text{for any } x, y$$

i.e.,  $\nabla f$  is Lipschitz continuous with constant  $L > 0$ .

### Theorem

*Gradient descent with fixed step size  $t \leq 1/L$  satisfies*

$$f(x^{(k)}) - f^* \leq \frac{1}{2tk} \|x^{(0)} - x^*\|_2^2$$

*and same result holds for backtracking with  $t$  replaced by  $\beta/L$ .*

We say gradient descent has convergence rate  $O(1/k)$ , i.e., it finds  $\varepsilon$ -suboptimal point in  $O(1/\varepsilon)$  iterations.

## Convergence analysis

Assume that  $f : \mathbb{R}^n \rightarrow \mathbb{R}$  convex and differentiable and additionally

$$\|\nabla f(x) - \nabla f(y)\|_2 \leq L \|x - y\|_2 \quad \text{for any } x, y$$

i.e.,  $\nabla f$  is Lipschitz continuous with constant  $L > 0$ .

### Theorem

Gradient descent with fixed step size  $t \leq 1/L$  satisfies

$$f(x^{(k)}) - f^* \leq \frac{1}{2tk} \|x^{(0)} - x^*\|_2^2$$

and same result holds for backtracking with  $t$  replaced by  $\beta/L$ .

We say gradient descent has convergence rate  $O(1/k)$ , i.e., it finds  $\varepsilon$ -suboptimal point in  $O(1/\varepsilon)$  iterations.

### Chứng minh.

Slide 20-25 in <http://www.seas.ucla.edu/~vandenbe/236C/lectures/gradient.pdf>



## Convergence under strong convexity

Reminder: **strong convexity** of  $f$  means  $f(x) - \frac{m}{2} \|x\|_2^2$  is convex for some  $m > 0$ .

Assuming Lipschitz gradient as before and also strong convexity:

### Theorem

*Gradient descent with fixed step size  $t \leq 2/(m + L)$  or with backtracking line search search satisfies*

$$f(x^{(k)}) - f^* \leq c^k \frac{L}{2} \|x^{(0)} - x^*\|_2^2$$

*where  $0 < c < 1$ .*

## Convergence under strong convexity

Reminder: **strong convexity** of  $f$  means  $f(x) - \frac{m}{2} \|x\|_2^2$  is convex for some  $m > 0$ .

Assuming Lipschitz gradient as before and also strong convexity:

### Theorem

*Gradient descent with fixed step size  $t \leq 2/(m + L)$  or with backtracking line search search satisfies*

$$f(x^{(k)}) - f^* \leq c^k \frac{L}{2} \|x^{(0)} - x^*\|_2^2$$

*where  $0 < c < 1$ .*

Rate under strong convexity is  $O(c^k)$ , exponentially fast, i.e., we find  $\varepsilon$ -suboptimal point in  $O(\log(1/\varepsilon))$  iterations.

# Convergence under strong convexity

Reminder: **strong convexity** of  $f$  means  $f(x) - \frac{m}{2} \|x\|_2^2$  is convex for some  $m > 0$ .

Assuming Lipschitz gradient as before and also strong convexity:

## Theorem

*Gradient descent with fixed step size  $t \leq 2/(m + L)$  or with backtracking line search search satisfies*

$$f(x^{(k)}) - f^* \leq c^k \frac{L}{2} \|x^{(0)} - x^*\|_2^2$$

where  $0 < c < 1$ .

Rate under strong convexity is  $O(c^k)$ , exponentially fast, i.e., we find  $\varepsilon$ -suboptimal point in  $O(\log(1/\varepsilon))$  iterations.


## Chứng minh.

Slide 26-27 in <http://www.seas.ucla.edu/~vandenbe/236C/lectures/gradient.pdf>



## Convergence rate

Called **linear convergence**,  
because looks linear on a  
semi-log plot.



(From B & V page 487)

Important note: contraction factor  $c$  in rate depends adversely on condition number  $L/m$ : higher condition number  $\Rightarrow$  slower rate.

Affects not only our upper bound... very apparent in practice too.

## A look at the conditions

A look at the conditions for a simple problem,  $f(\beta) = \frac{1}{2} \|y - X\beta\|_2^2$ .

Lipschitz continuity of  $\nabla f$ :

- ▶ This mean  $\nabla^2 f(\beta) \preceq L I$ .
- ▶ As  $\nabla^2 f(\beta) = X^T X$ , we have  $L = \sigma_{\max}(X^T X)$ .

Strong convexity of  $f$ :

- ▶ This mean  $\nabla^2 f(\beta) \succeq m I$ .
- ▶ As  $\nabla^2 f(\beta) = X^T X$ , we have  $m = \sigma_{\min}(X^T X)$ .
- ▶ If  $X$  is wide (i.e.,  $X$  is  $n \times p$  with  $p > n$ ), then  $\sigma_{\min}(X^T X) = 0$ , and  $f$  can't be strongly convex.
- ▶ Even if  $\sigma_{\min}(X^T X) > 0$ , can have a very large condition number  $L/m = \sigma_{\max}(X^T X)/\sigma_{\min}(X^T X)$ .

## Practicalities

Stopping rule: stop when  $\|\nabla f(x)\|_2$  is small

- ▶ Recall  $\nabla f(x^*) = 0$  at solution  $x^*$
- ▶ If  $f$  is strongly convex with parameter  $m$ , then

$$\|\nabla f(x)\|_2 \leq \sqrt{2m\varepsilon} \implies f(x) - f^* \leq \varepsilon.$$

## Practicalities

Stopping rule: stop when  $\|\nabla f(x)\|_2$  is small

- ▶ Recall  $\nabla f(x^*) = 0$  at solution  $x^*$
- ▶ If  $f$  is strongly convex with parameter  $m$ , then

$$\|\nabla f(x)\|_2 \leq \sqrt{2m\varepsilon} \implies f(x) - f^* \leq \varepsilon.$$

Pros and cons of gradient descent:

- ▶ Pro: simple idea, and each iteration is cheap (usually)
- ▶ Pro: fast for well-conditioned, strongly convex problems
- ▶ Con: can often be slow, because many interesting problems aren't strongly convex or well-conditioned
- ▶ Con: can't handle nondifferentiable functions.

## Can we do better?

Gradient descent has  $O(1/\varepsilon)$  convergence rate over problem class of convex, differentiable functions with Lipschitz gradients.

**First-order method:** iterative method, which updates  $x^{(k)}$  in

$$x^{(0)} + \text{span}\{\nabla f(x^{(0)}), \nabla f(x^{(1)}), \dots, \nabla f(x^{(k-1)})\}.$$

## Can we do better?

Gradient descent has  $O(1/\varepsilon)$  convergence rate over problem class of convex, differentiable functions with Lipschitz gradients.

**First-order method:** iterative method, which updates  $x^{(k)}$  in

$$x^{(0)} + \text{span}\{\nabla f(x^{(0)}), \nabla f(x^{(1)}), \dots, \nabla f(x^{(k-1)})\}.$$

### Theorem (Nesterov)

For any  $k \leq (n - 1)/2$  and any starting point  $x^{(0)}$ , there is a function  $f$  in the problem class such that any first-order method satisfies

$$f(x^{(k)}) - f^* \geq \frac{3L \|x^{(0)} - x^*\|_2^2}{32(k+1)^2}.$$

Can attain rate  $O(1/k^2)$ , or  $O(1/\sqrt{\varepsilon})$ ? Answer: yes (we'll see)!

## What about nonconvex functions?

Assume  $f$  is differentiable with Lipschitz gradient as before, but now **nonconvex**. Asking for optimality is too much. So we'll settle for  $x$  such that  $\|\nabla f(x)\|_2 \leq \varepsilon$ , called  **$\varepsilon$ -stationarity**.

---

<sup>1</sup>Carmon et al. (2017), "Lower bounds for finding stationary points I"

## What about nonconvex functions?

Assume  $f$  is differentiable with Lipschitz gradient as before, but now **nonconvex**. Asking for optimality is too much. So we'll settle for  $x$  such that  $\|\nabla f(x)\|_2 \leq \varepsilon$ , called  **$\varepsilon$ -stationarity**.

### Theorem

Gradient descent with fixed step size  $t \leq 1/L$  satisfies

$$\min_{i=0,\dots,k} \|\nabla f(x^{(i)})\|_2 \leq \sqrt{\frac{2(f(x^0) - f^*)}{t(k+1)}}.$$

Thus gradient descent has rate  $O(1/\sqrt{k})$ , or  $O(1/\varepsilon^2)$ , even in the nonconvex case for finding stationary points.

This rate **cannot be improved** (over class of differentiable functions with Lipschitz gradients) by any deterministic algorithm<sup>1</sup>.

<sup>1</sup>Carmon et al. (2017), "Lower bounds for finding stationary points I"

## Proof

Key steps:

- $\nabla f$  Lipschitz with constant  $L$  means

$$f(y) \leq f(x) + \nabla f(x)^T(y - x) + \frac{L}{2} \|y - x\|_2^2, \quad \forall x, y.$$

- Plugging in  $y = x^+ = x - t\nabla f(x)$ ,

$$f(x^+) \leq f(x) - \left(1 - \frac{Lt}{2}\right) t \|\nabla f(x)\|_2^2.$$

- Taking  $0 < t \leq 1/L$ , and rearranging,

$$\|\nabla f(x)\|_2^2 \leq \frac{2}{t}(f(x) - f(x^+)).$$

- Summing over iterations

$$\sum_{i=0}^k \|\nabla f(x^{(i)})\|_2^2 \leq \frac{2}{t}(f(x^{(0)}) - f(x^{(k+1)})) \leq \frac{2}{t}(f(x^{(0)}) - f^*).$$

- Lower bound sum by  $(k+1) \min_{i=0,1,\dots} \|\nabla f(x^{(i)})\|_2^2$ , conclude.

## References and further reading

-  S. Boyd and L. Vandenberghe (2004), *Convex optimization*, Chapter 9
-  T. Hastie, R. Tibshirani and J. Friedman (2009), *The elements of statistical learning*, Chapters 10 and 16
-  Y. Nesterov (1998), *Introductory lectures on convex optimization: a basic course*, Chapter 1
-  L. Vandenberghe, *Lecture notes for EE 236C*, UCLA, Spring 2011-2012